

Report

Subject: Plan of Execution of Queries.

Students: CHARFAOUI Younes, BOURBAI Ismail.

Github Repository: https://github.com/IsmailBourbie/master-one-practical-work/tree/master/db_dm/PW03

Step One:

We've installed the oracle 11g desktop class in our computers, and we've used Datagrip IDE for manipulating and queries the DBMS

Step Two:

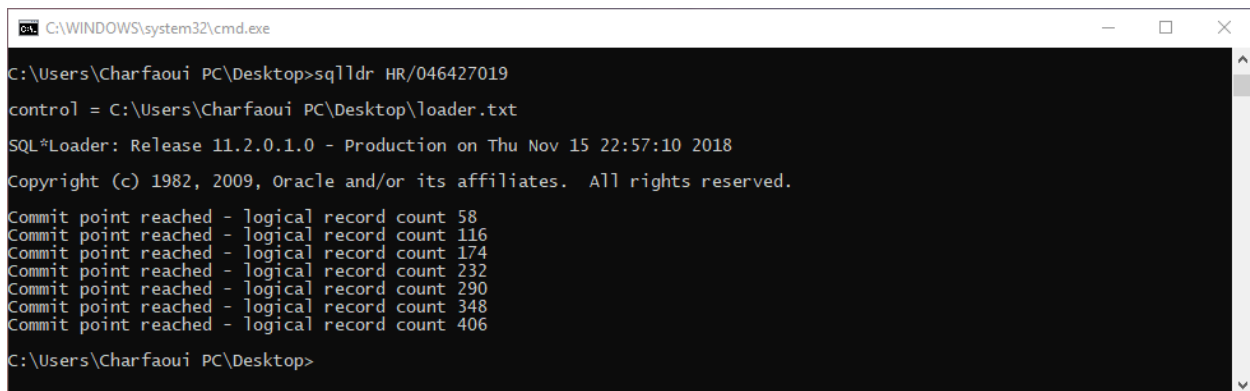
In this step we create the table provided in the practical work sheet .Let's Execute the Following Line of SQL Script, the other queries are in the github repos.

```
CREATE TABLE LINEORDER (  
    LO_ORDERKEY          INTEGER,  
    LO_LINENUMBER        INTEGER,  
    LO_CUSTKEY            INTEGER    NOT NULL,  
    LO_PARTKEY            INTEGER    NOT NULL,  
    LO_SUPPKEY            INTEGER    NOT NULL,  
    LO_ORDERDATE          INTEGER    NOT NULL,  
    LO_ORDERPRIOTITY      VARCHAR(15) NOT NULL,  
    LO_SHIPPRIOTITY        INTEGER,  
    LO_QUANTITY            INTEGER,  
    LO_EXTENDEDPRICE      INTEGER,  
    LO_ORDTOTALPRICE      INTEGER,  
    LO_DISCOUNT          INTEGER,  
    LO_REVENUE             INTEGER,  
    LO_SUPPLYCOST          INTEGER,  
    LO_TAX                 INTEGER,  
    LO_COMMITDATE          INTEGER    NOT NULL,  
    LO_SHIPMODE            VARCHAR(10) NOT NULL  
);
```

Figure 1 Example of Create Query

Step Three:

In This Step we have used the SQL*Loader to load our prepared data into the tables, the following screen shot demonstrate one of the tables that:



```
C:\WINDOWS\system32\cmd.exe
C:\Users\Charfaoui PC\Desktop>sqlldr HR/046427019
control = C:\Users\Charfaoui PC\Desktop\loader.txt
SQL*Loader: Release 11.2.0.1.0 - Production on Thu Nov 15 22:57:10 2018
Copyright (c) 1982, 2009, Oracle and/or its affiliates. All rights reserved.

Commit point reached - logical record count 58
Commit point reached - logical record count 116
Commit point reached - logical record count 174
Commit point reached - logical record count 232
Commit point reached - logical record count 290
Commit point reached - logical record count 348
Commit point reached - logical record count 406
C:\Users\Charfaoui PC\Desktop>
```

Figure 2 SQL Loader Demonstration

In The Loader file we specify the table and from which file we are taking the data, here is an model:

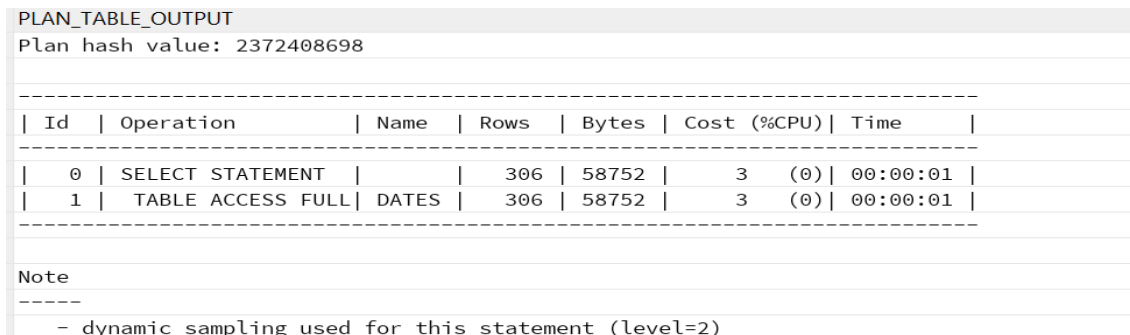
LOAD DATA INFILE "path_of_datafile" INTO TABLE "name_of_table" FIELDS TERMINATED BY separator (columns)

The other file are in the repo: https://github.com/IsmailBourbie/master-one-practical-work/tree/master/db_dm/PW03/loader_files

Step Four:

For a DBA there is very handy tool called Explain Plan, and this one help the DBA for tracking how the query is executed to help it debug and find an optimal way to reduce the time and make accurate result, in this Practical Work we explored different queries and its plan, Here is the Query and the corresponding plan:

EXPLAIN PLAN FOR SELECT * FROM DATES;



Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		306	58752	3 (0)	00:00:01
1	TABLE ACCESS FULL	DATES	306	58752	3 (0)	00:00:01

Note

- dynamic sampling used for this statement (level=2)

Figure 3 Result Plan

We made another test with a complex query, here is both of them:

```

EXPLAIN PLAN FOR
select
    sum(lo_extendedprice*lo_discount) as revenue
from
    lineorder, dates
where
    lo_orderdate = d_datekey
    and d_yearmonthnum = 199401
    and lo_discount between 4 and 6
    and lo_quantity between 26 and 35;

```

Figure 4 Complex query for The EXPLAIN PLAN

PLAN_TABLE_OUTPUT									
Plan hash value: 2553672856									

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time			

0	SELECT STATEMENT		1	78	19822 (2)	00:03:58			
1	SORT AGGREGATE		1	78					
* 2	HASH JOIN		1	78	19822 (2)	00:03:58			
* 3	TABLE ACCESS FULL	DATES	1	26	3 (0)	00:00:01			
* 4	TABLE ACCESS FULL	LINEORDER	338K	16M	19816 (2)	00:03:58			

Predicate Information (identified by operation id):									

2 - access("LO_ORDERDATE"="D_DATEKEY")									
3 - filter("D_YEARMONTHNUM"=199401)									
4 - filter("LO_DISCOUNT">=4 AND "LO_DISCOUNT"<=6 AND									
"LO_QUANTITY">=26 AND "LO_QUANTITY"<=35)									
Note									

- dynamic sampling used for this statement (level=2)									

Figure 5 Plan of Execution of Last Query

The Other Examples are in the repository. https://github.com/IsmailBourbie/master-one-practical-work/tree/master/db_dm/PW03/screen_shots

Step Five:

At this point we know how to see the way the query is executing, but how tune some parameter? The way to do that is throughout using hints, hints are a way for telling the DBMS how to execute queries, we 'have used The Nested Loop Hint instead of the default Hash Join to Show the difference, here is the query and the associated plan.

```
EXPLAIN PLAN FOR
select /*+ use_nl(lineorder,dates,part,supplier) */
      sum(lo_revenue), d_year, p_brand
from
      lineorder, dates, part, supplier
where
      lo_orderdate = d_datekey
      and lo_partkey = p_partkey
      and lo_suppkey = s_suppkey
      and p_category = 'MFGR#12'
      and s_region = 'AMERICA'
group by
      d_year, p_brand
order by
      d_year, p_brand;
```

Figure 6 Query Use A Hint to optimize

PLAN_TABLE_OUTPUT								
Plan hash value: 1339098462								

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time		

0	SELECT STATEMENT		13783	1763K	943K (1)	03:08:48		
1	SORT GROUP BY		13783	1763K	943K (1)	03:08:48		
2	NESTED LOOPS		13783	1763K	943K (1)	03:08:48		
3	NESTED LOOPS		72923	7833K	386K (1)	01:17:23		
* 4	HASH JOIN		269K	21M	20536 (2)	00:04:07		
5	VIEW	VW_GBF_13	8094	252K	688 (1)	00:00:09		
6	HASH GROUP BY		8094	189K	688 (1)	00:00:09		
* 7	TABLE ACCESS FULL	PART	8094	189K	687 (1)	00:00:09		
8	TABLE ACCESS FULL	LINEORDER	6494K	322M	19802 (2)	00:03:58		
* 9	TABLE ACCESS FULL	DATES	1	26	1 (0)	00:00:01		
* 10	TABLE ACCESS FULL	SUPPLIER	1	21	8 (0)	00:00:01		

Predicate Information (identified by operation id):								

4 - access("LO_PARTKEY"="ITEM_1")								
7 - filter("P_CATEGORY"='MFGR#12')								
9 - filter("LO_ORDERDATE"="D_DATEKEY")								
10 - filter("S_REGION"='AMERICA' AND "LO_SUPPKEY"="S_SUPPKEY")								
Note								

- dynamic sampling used for this statement (level=2)								

Figure 7 the Plan of Nested Loop Example

In The Operation with id 3 and 2, we see that the DBMS has used Nested Loop in The Plan of executing the query and this demonstrate how we use hints.

Step Six:

In This last step, our goal is to use an optimization technique to make the query work faster as possible for this we used the indexes and horizontal fragmentation, the queries are in the repository and one of the testing query and it's correspond plan are below :

Other Queries: https://github.com/IsmailBourbie/master-one-practical-work/blob/master/db_dm/PW03/queries.sql

```

EXPLAIN PLAN FOR
select
    c_city, s_city, d_year, sum(lo_revenue) as revenue
from
    customer, lineorder, supplier, dates
where
    lo_custkey = c_custkey
    and lo_suppkey = s_suppkey
    and lo_orderdate = d_datekey
    and (c_city='UNITED KI1' or c_city='UNITED KI5')
    and (s_city='UNITED KI1' or s_city='UNITED KI5')
    and d_yearmonth = 'Dec1997'
group by
    c_city, s_city, d_year
order by
    d_year asc, revenue desc;;

```

Figure 8 Query to demonstrate Index Optimization

PLAN_TABLE_OUTPUT							
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
0	SELECT STATEMENT		1	123	132 (2)	00:00:02	
1	SORT ORDER BY		1	123	132 (2)	00:00:02	
2	HASH GROUP BY		1	123	132 (2)	00:00:02	
3	NESTED LOOPS						
4	NESTED LOOPS		1	123	130 (0)	00:00:02	
5	NESTED LOOPS		1	103	23 (0)	00:00:01	
6	MERGE JOIN CARTESIAN		1	51	8 (0)	00:00:01	
7	INLIST ITERATOR						
8	TABLE ACCESS BY INDEX ROWID	SUPPLIER	19	380	8 (0)	00:00:01	
* 9	INDEX RANGE SCAN	S_CITY_IDX	8		2 (0)	00:00:01	
10	BUFFER SORT		1	31	0 (0)	00:00:01	
11	TABLE ACCESS BY INDEX ROWID	DATES	1	31	0 (0)	00:00:01	
* 12	INDEX RANGE SCAN	D_YEARMONTH_IDX	1		0 (0)	00:00:01	
* 13	TABLE ACCESS BY INDEX ROWID	LINEORDER	11	572	15 (0)	00:00:01	
* 14	INDEX RANGE SCAN	LO_ORDERDATE_IDX	35		6 (0)	00:00:01	
15	INLIST ITERATOR						
* 16	INDEX RANGE SCAN	C_CITY_IDX	116		2 (0)	00:00:01	
* 17	TABLE ACCESS BY INDEX ROWID	CUSTOMER	1	20	107 (0)	00:00:02	

Predicate Information (identified by operation id):							

9 - access("S_CITY"='UNITED KI1' OR "S_CITY"='UNITED KI5')							
12 - access("D_YEARMONTH"='Dec1997')							
13 - filter("LO_SUPPKEY"="S_SUPPKEY")							
14 - access("LO_ORDERDATE"="D_DATEKEY")							
16 - access("C_CITY"='UNITED KI1' OR "C_CITY"='UNITED KI5')							

Figure 9 Plan after Index Modification

Conclusion:

In This Practical word we saw that oracle DBMS is a great one , its offer loading data from various input , help with monitoring the query plan execution and also give the database administrator the tools and ways to optimize the database structure and the corresponding query via multiple parameters, another future which is a great about it but really made us crazy about oracle is the security provided , really in the first time your hand will be very dirty just to enter into the DBMS and start making some queries, really it was very difficult to just enter into it, lastly I would really recommend this DBMS for A very large enterprise that need security and performance.